



TRƯỜNG ĐẠI HỌC
BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY
OF SCIENCE AND TECHNOLOGY

Flight Delays and Cancellations Pipeline

Group 1

Phạm Anh Tú 20225463

Etienne Fontbonne 20250006S

Elliott Moskowicz 20250075S

Florian Dorchies 20250072S

Philippe Marchand 20250083S

ONE LOVE. ONE FUTURE.

Problem Statement

Air travel delays are a persistent issue affecting passengers, airlines, and airport operations worldwide.

Delays can arise from a complex interplay of factors such as :

- weather conditions
- air traffic congestion
- carrier performance
- airport infrastructure
- scheduling inefficiencies.

We aim to **process** and **visualize** various statistics about this dataset to inform customers about the performance of different airlines and routes.

2015 Flight Delays and Cancellations dataset


Tracks the on-time performance of domestic flights in the USA operated by large air carriers.

- 3 tables:
 - airlines : 2 columns
 - airports : 7 columns
 - flights : 31 columns
- 4 data types
- Source : kaggle

Over 1M rows in the dataset.

Proposed pipeline

Kappa architecture



Kafka

- **Main role:** Data ingestion
- **Producer :** sends flights events into a Kafka topic
- **Broker :** central server that stores and distributes messages
- **Using Zookeeper :**
 - Zookeeper ensures consistency across the cluster.
 - Handles broker leader election and monitors node status
 - Produces through a python file

Spark

- **Main role:** Data Processing
- **Transformations:**
 - Parse JSON messages using robust type casting
 - Read and cached static data
 - Handle missing values and data formatting
 - Add timestamp
 - Custom UDF: separate on-time flights and delayed flights according to US laws.

Spark

- **Aggregations:**
 - Airline-level statistics (on-time, delayed, cancelled flights and calculate average delay).
 - Route-level average delays enriched with geographic information.
 - Hourly delay trends (avg, std, median, ...) and statistical summaries.
 - Delay statistics grouped by delay reasons

Spark

- **Key optimizations:**
 - Controlled micro-batch size and backpressure
 - Partition alignment with Kafka
 - Caching static datasets
 - Resource limits via Docker/Kubernetes
- **Result:**
 - Stable startup
 - Predictable latency


Cassandra

- **Role:** Serving database for Grafana
- Stores 4 tables:
 - `airline_stats`: Stores airline-level performance metrics aggregated over streaming batches
 - `delay_by_reason`: Stores aggregated delay statistics grouped by delay category
 - `route_stats`: Stores route-level delay analytics between origin and destination airports.
 - `hourly_stats`: Stores time-series statistics aggregated by scheduled departure hour

Components

Grafana

- Role: End of pipeline visualization




HDFS

- **Main role :** Cold data storage and checkpoints
- **Data archival :**
 - Save of enriched flights for archival purposes
- **Checkpoints :**
 - Acts as a safe place for Spark streaming checkpoint
 - This allows recovery of streaming jobs after failures, state management and improved resistance

From Docker to Kubernetes

- Firstly started on Docker for easier development
- Then switch to Kubernetes, using Minikube
- Problematic :
 - Need to convert `docker-compose.yml` to Kubernetes manifest
 - Preserve process starting order




From Docker to Kubernetes

- Solution :
 - Use of InitContainers to preserve startup dependencies
 - Define StatefulSet for stable containers (Kafka, Cassandra, ...)
 - Define Deployment to enable easier scaling and updates (Spark, Grafana, ...)
 - Replace configuration files by ConfigMaps
 - Define PVC to ensure durable storage
 - Automatic deployment using a bash file

From Docker to Kubernetes

- Issues :
 - HDFS :
 - Not resistant to pressure
 - Needed too much resources to be stable
 - -> could not be used in the **Kubernetes** deployment
 - Consequences :
 - Local checkpoint in Spark processing
 - No cold storage
 - No access to local repository :
 - Solution :
 - Uses of mounted ConfigMap during deployment
 - Creation of a custom Docker Image for Kafka



From Docker to Kubernetes

- Deployment script :
 - Checks minikube and dataset status
 - Applying namespace and then all manifests
 - Builds the Kafka Docker Image
 - Loads the ConfigMap
- Then validation with :
 - `kubectl get pods`
 - `kubectl logs <pod>`
 - `minikube logs`

Challenges

- Not all images are free : Bitnami Spark, Grafana MongoDB plugins .
 - => caused changes in the pipeline
- Latency between data generation and availability for processing
 - Controlling micro-batches size, flow control
- Spark streaming from Kafka produced null metrics and startup delays due to parsing and backlog issues
 - Ingest fields as StringType
 - Data quality checks



HUST

THANK YOU !